Asimiyu Musa

Data Engineering & Data Science • Client Portfolio

I build scalable, compliant, and insight-driven data platforms across governance, migration, batch/streaming pipelines, and analytics.

Focus areas

- Enterprise data governance (NDPA / GDPR-aligned)
- Cloud & hybrid migration (AWS / Azure / GCP)
- Batch + real-time pipelines (Airflow, Kafka)
- Analytics platforms (PostgreSQL, BI, NLP)

What I deliver

Engagements designed to reduce risk, improve reliability, and speed up insight.

Data Platform Architecture

Modern data platform design, data models, and scalable patterns for growth.

Governance & Compliance

Policies, classification, RBAC, auditing, and control frameworks aligned to NDPA.

Cloud Migration

Phased migrations with minimal downtime; cost/performance optimisation.

Pipelines & Orchestration

Batch and streaming pipelines with quality checks and observability.

Analytics & BI

Decision support dashboards and products powered by clean, reliable datasets.

Case Study: NDPA Data Governance (Utility Sector)

A compliance-ready governance framework for sensitive customer & operational data.

Problem

Utility data is highly sensitive, but governance is often fragmented and not aligned to NDPA.

Approach

- Classification & handling rules
- RBAC access controls
- Audit logging & traceability
- Risk assessment and mitigation
- Compliance workflows

Outcome

Improved audit readiness, reduced data-risk, and created a scalable governance model.

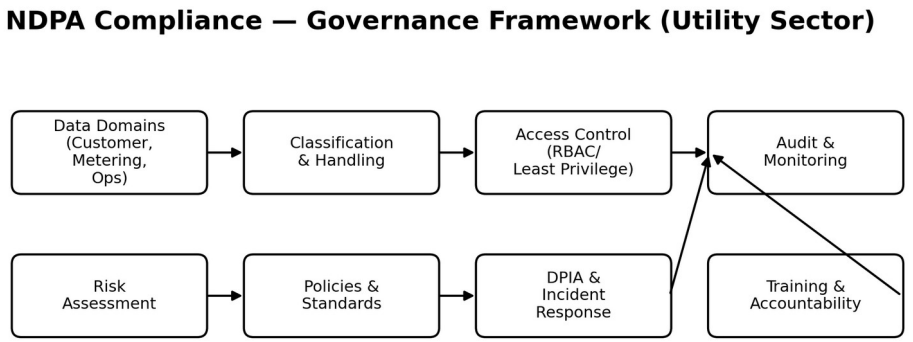

Case Study: Batch Migration to Cloud

Phased migration from on-prem data warehouse to cloud with minimal downtime.

Highlights

- Phased cutover plan
- Cost/performance optimisation
- Data validation & reconciliation
- Rollback and continuity controls

Outcome

Improved scalability and resilience with a cloud-ready analytics foundation.

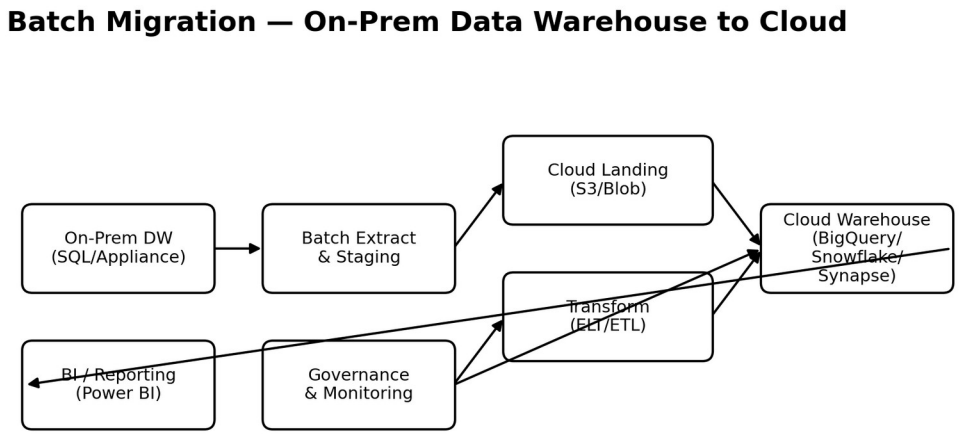

Case Study: Auto-Intel Platform

Automated data intelligence platform integrating scraping, orchestration, APIs, NLP and BI.

What it does

- Web scraping + extraction
- Validation + transformation
- Airflow orchestration
- PostgreSQL storage
- API layer
- NLP-powered insight + BI

Outcome

End-to-end automation of market intelligence workflows with faster, more reliable insights.

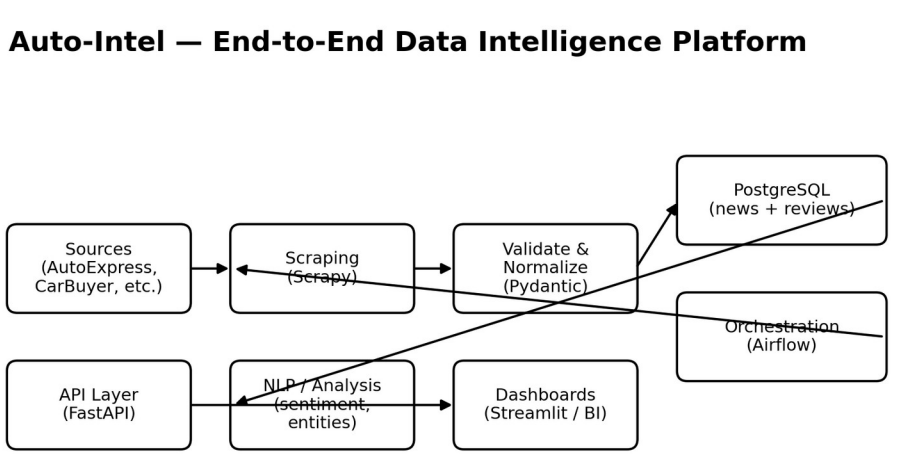

Case Study: Real-Time Streaming Pipeline

Kafka + Airflow pipeline for near real-time ingestion and analytics readiness.

Highlights

- Event-driven ingestion
- Stream processing + validation
- Orchestration and monitoring
- Analytics-ready outputs

Outcome

Near real-time reporting with improved operational responsiveness.

Streaming Pipeline — Kafka + Airflow

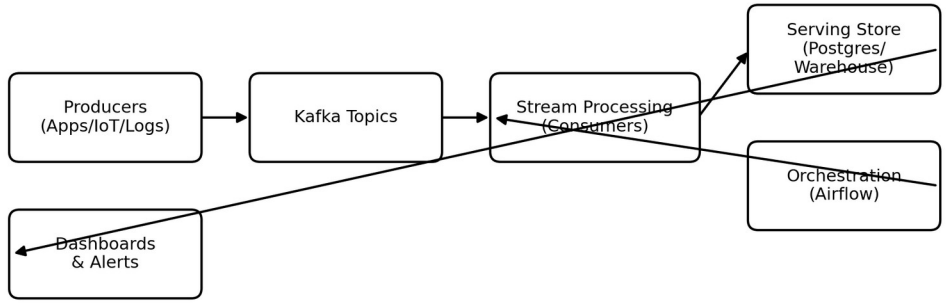

Additional projects

More examples of end-to-end delivery across domains.

Enterprise Big Data QA Pipeline

Profiling, automated checks, and validation to improve trust in analytics.

Real Estate ETL (PostgreSQL)

Standardised property records; improved querying and reporting performance.

Retail Sales Integration (Airbyte + S3)

Consolidated multi-source sales data into an analytics-ready data lake.

Rural Water Analytics

Multi-source analysis and roadmap for infrastructure and sustainability planning.

Technology stack

Tools I use to build reliable data products.

Languages

Python, SQL

Data Stores

PostgreSQL, BigQuery

Orchestration

Apache Airflow

Streaming

Apache Kafka

Cloud

AWS, Azure, GCP

BI & Apps

Streamlit, Power BI

Governance

NDPA-aligned frameworks

Engagement models & contact

Let’s discuss your data platform goals and choose the right approach.

Engagement options

- One-off delivery (project-based)
- Advisory / architecture review
- Retainer (ongoing improvements)
- Contract role (short- or mid-term)

What you get

Clear deliverables, documentation, quality controls, and knowledge transfer.

Contact

GitHub Pages website
(asimiyu-musa.github.io)

GitHub
(github.com/asimiyu-musa)

LinkedIn
(add your link)

Next step: Share your goal + data sources, and I’ll propose an approach and deliverables.